


NLP and the Web – WS 2024/2025



TECHNISCHE
UNIVERSITÄT
DARMSTADT

Lecture 3 Information Retrieval I

Dr. Thomas Arnold
Hovhannes Tamoyan
Kexin Wang



Ubiquitous Knowledge Processing Lab
Technische Universität Darmstadt

Syllabus (tentative)

<u>Nr.</u>	<u>Lecture</u>
01	Introduction / NLP basics
02	Foundations of Text Classification
03	IR – Introduction, Evaluation
04	IR – Word Representation, Data Collection
05	IR – Re-Ranking Methods
06	IR – Language Domain Shifts, Dense / Sparse Retrieval
07	LLM – Language Modeling Foundations
08	LLM – Neural LLM, Tokenization
09	LLM – Transformers, Self-Attention
10	LLM – Adaption, LoRa, Prompting
11	LLM – Alignment, Instruction Tuning
12	LLM – Long Contexts, RAG
13	LLM – Scaling, Computation Cost
14	Review & Preparation for the Exam

Today

IR – Introduction, Evaluation

1 Introduction

- Inverted Index
- Search & Relevance
- TF-IDF & BM25

2 Evaluation

- Precision & Recall
- MRR & MAP
- nDCG

Information Retrieval

“Elephant weight”



How
Relevant?

Document






Information Retrieval (finding the needle in the haystack)

“Elephant weight”



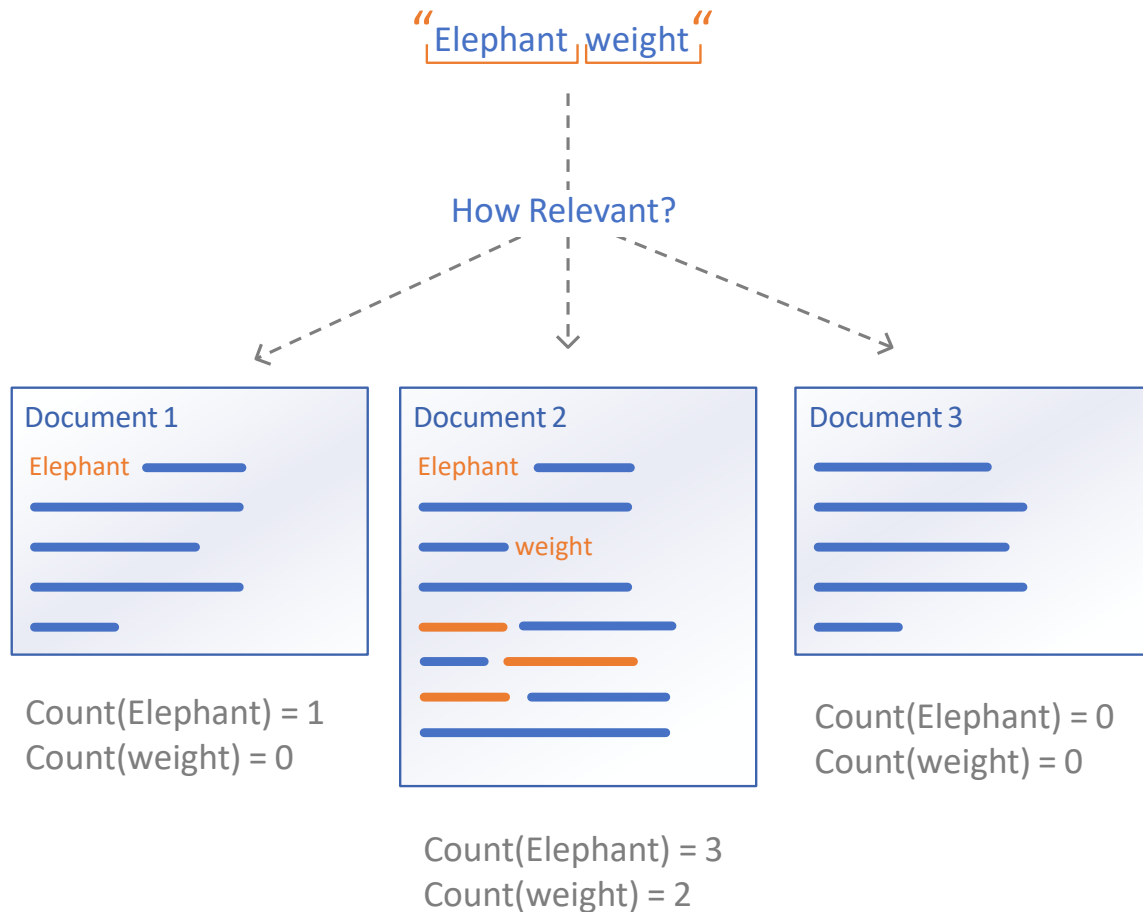
**How
Relevant?**



Notes on terminology

- **Documents** can be anything: a web page, word file, text file, article ...
(we assume it to be text for the moment)
 - A lot of details to look out for: encoding, language, hierarchy, fields, ...
- **Collection:** A set of documents (we assume it to be static for the moment)
- **Relevance:** Does a document satisfy the information need of the user and does it help complete the user's task?

Relevance (based on text content)



- If a word appears more often -> more relevant
- Solution: count the words
- If a document is longer, words will tend to appear more often -> take into account the document length
- Counting only when we have a query is inefficient

Today

IR – Introduction, Evaluation

1 Introduction

- **Inverted Index**
- Search & Relevance
- TF-IDF & BM25

2 Evaluation

- Precision & Recall
- MRR & MAP
- nDCG

Inverted Index

- Inverted index allows to efficiently retrieve documents from large collections
- Inverted index stores all statistics per term (that the scoring model needs)
 - **Document frequency:** how many documents contain the term
 - **Term frequency per document:** how often does the term appear per document
 - Document length
 - Average document length
- Save statistics in a format that is accessible **by a given term**
- Save metadata of a document (Name, location of the full text, etc..)

Inverted Index

Document data

Document Ids & Metadata:

```
[0] = ("Wildlife", "location",...)  
[1] = ("Zoo Vienna" ,...)  
...
```

Document Lengths:

```
[0] = 231 [1] = 381 ...
```

Term data

"elephant" =>

1:5	2:1	3:5	4:5	...
-----	-----	-----	-----	-----

"lion" =>

1:2	7:1	9:2	...
-----	-----	-----	-----

"weight" =>

4:1	6:4	...
-----	-----	-----

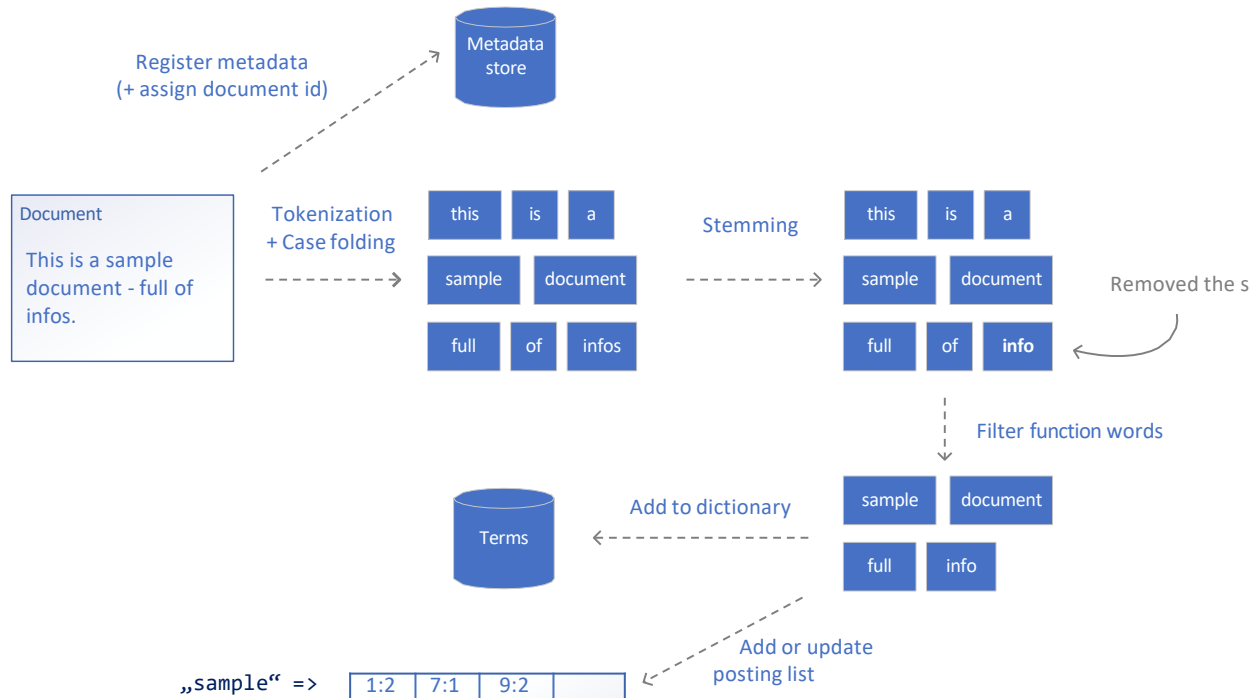
...

DocId

Term Frequency

- Every document gets an internal document id
- Term dictionary is saved as a search friendly data structure (more on that later)
- Term Frequencies are stored in a "posting list" = a list of doc id, frequency pairs

Creating the Inverted Index



- Simplified example pipeline
- Linguistic models are language dependent
- A query text and a document text both have to undergo the same steps

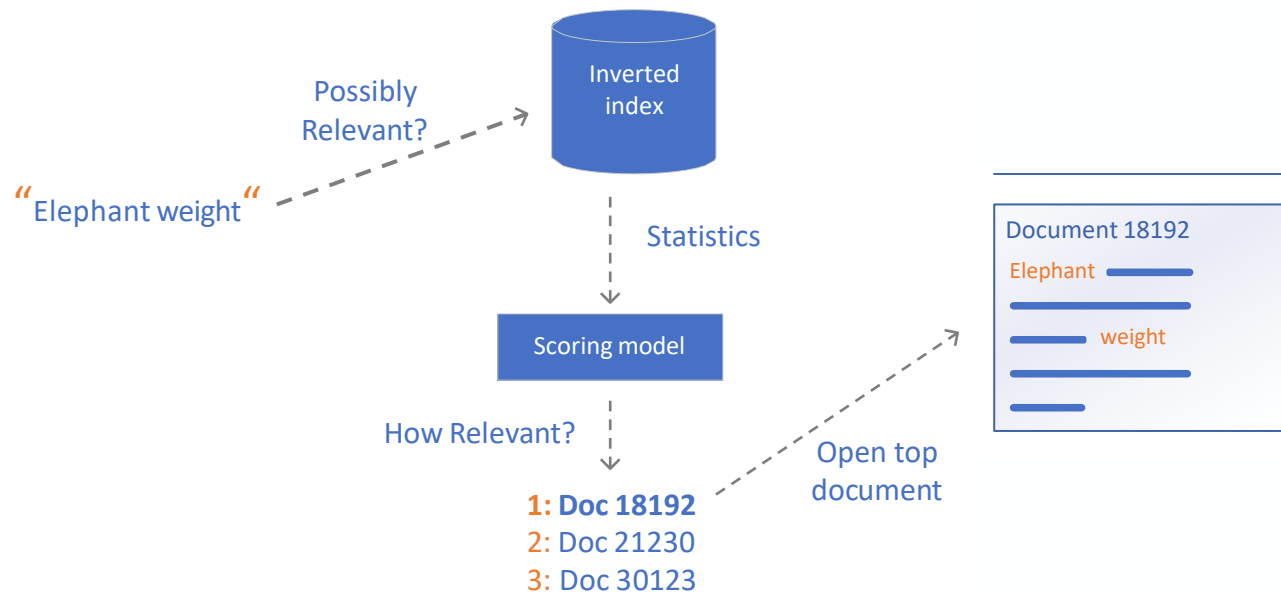
Today

IR – Introduction, Evaluation

- 1 Introduction
 - Inverted Index
 - **Search & Relevance**
 - TF-IDF & BM25

- 2 Evaluation
 - Precision & Recall
 - MRR & MAP
 - nDCG

Querying the Inverted Index



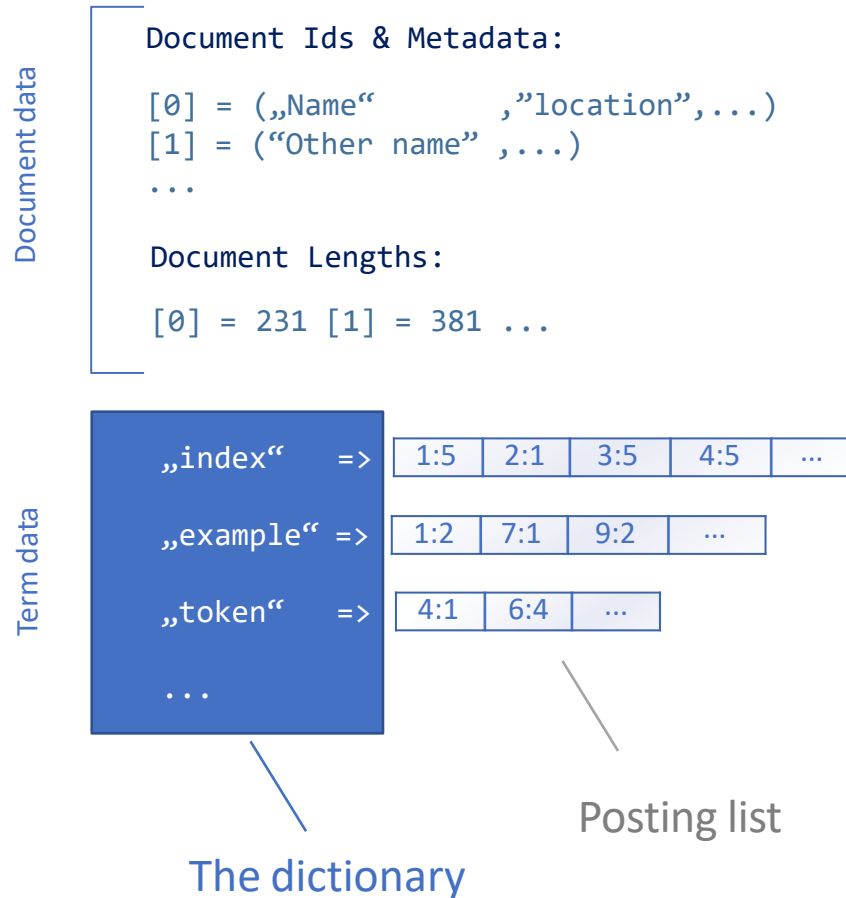
- No need to read full documents
- Only operate on frequency numbers of potentially relevant documents*
- Sort documents based on relevance score – retrieve most relevant documents

* it's not that easy because a document could be relevant without containing the exact query terms – but for now keep it simple

Types of queries (including, but not limited to)

- **Exact matching:** match full words and concatenate multiple query words with “or”
- **Boolean queries:** “and” / “or” / “not” operators between words
- **Expanded queries:** automatically incorporate synonyms and other similar or relevant words into the query
- **Wildcard queries, phrase queries, phonetic queries** (e.g. Soundex) ...

Inverted Index: Dictionary



- Dictionary<T> maps text to T
 - T is a posting list or potentially other data about the term depending on the index
- Wanted properties:
 - Random lookup
 - Fast (creation & especially lookup)
 - Memory efficient (keep the complete dictionary in memory)
- Naturally, there are a lot of choices

Scoring model

- Input: statistics, Output: floating point value (i.e. the score)
- Evaluated pairwise – 1 query, 1 document: $score(q, d)$
- Capture the notion of relevance in a mathematical model

Today we focus on free-text queries & „ad-hoc“ document retrieval

(document content only)

Search algorithm

float *Scores*={}

for each query term q

 fetch posting list for q

for each pair($d, tf_{t,d}$) in posting list

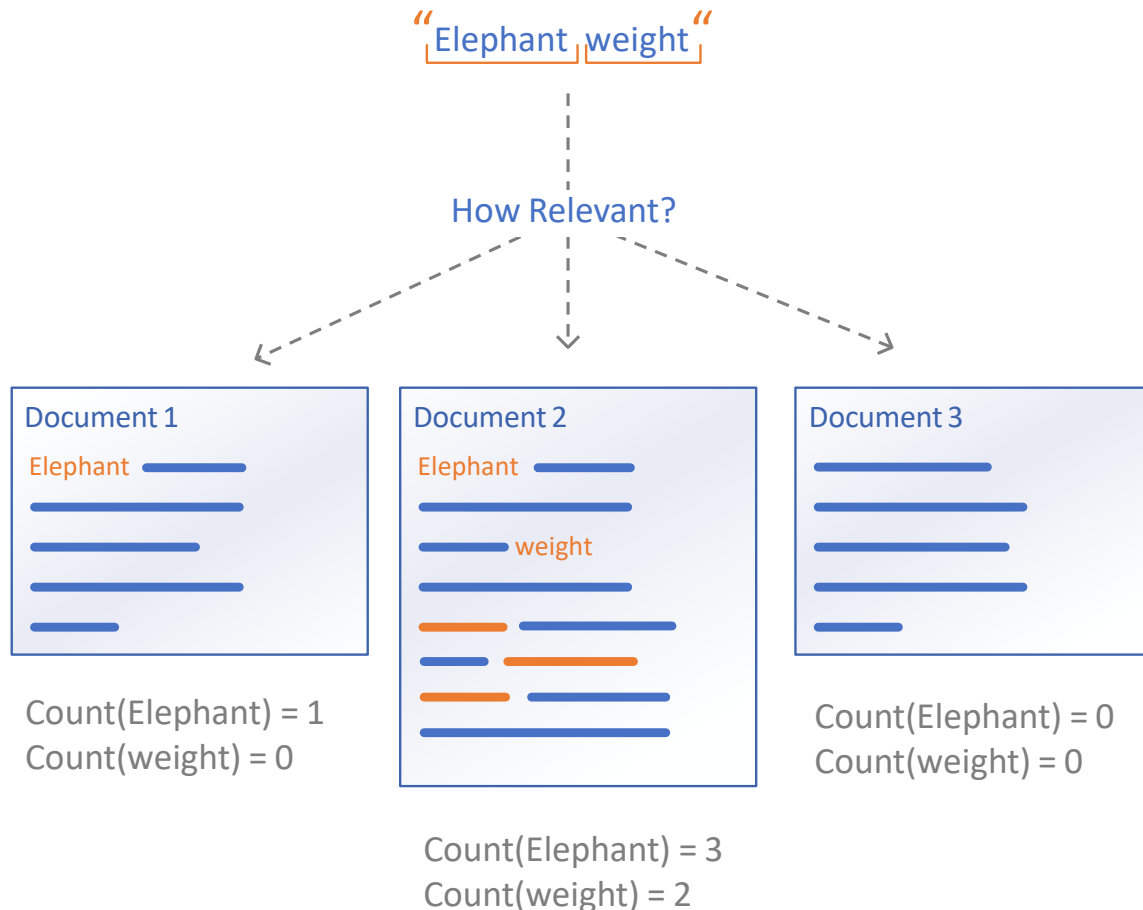
if d not in *Scores* **do** *Scores*[d]=0

Scores[d] += score($q, d, tf_{t,d}, \dots$)

return Top K entries of *Scores*

We transform information back to a document centric view (from the term centric view in the inverted index)

Relevance




- If a word appears more often → more relevant
- Solution: **count the words**
- If a document is longer, words will tend to appear more often → take into account the document length

Relevance limitations

- “Relevance” means relevance to the need rather than to the query
 - “Query” is shorthand for an instance of information need, its initial verbalized presentation by the user
- Relevance is assumed to be a binary attribute
 - A document is either relevant to a query/need or it is not
- We need these oversimplifications to create & evaluate mathematical models

From: A probabilistic model of information retrieval: development and comparative experiments, Spärck Jones et al. <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.134.6108&rep=rep1&type=pdf>



Today

IR – Introduction, Evaluation

1 Introduction

- Inverted Index
- Search & Relevance
- **TF-IDF & BM25**

2 Evaluation

- Precision & Recall
- MRR & MAP
- nDCG

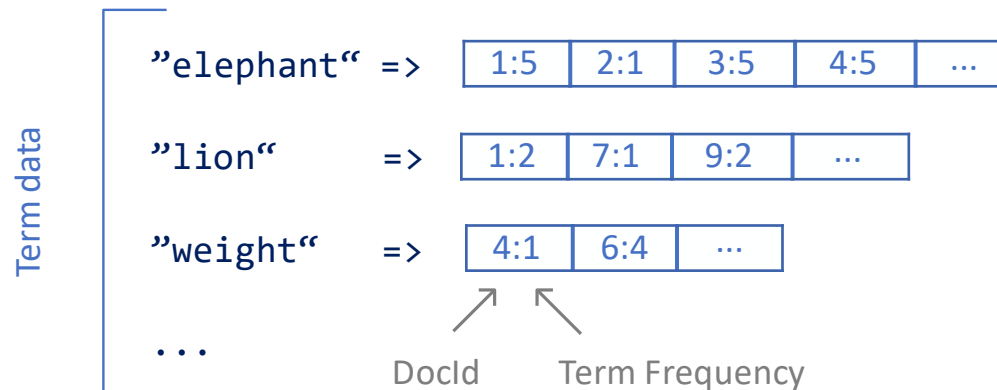
Term Frequency – conceptional data view

- **Bag of words:** word order is not important
- First step for a retrieval model: number of occurrences counts!
- $tf_{t,d}$ number of occurrences of term t in document d

		Documents					
		Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Terms	Antony	157	73	0	0	0	0
	Brutus	4	157	0	1	0	0
	Caesar	231	227	0	2	1	1
	Calpurnia	0	10	0	0	0	0
	Cleopatra	57	0	0	0	0	0
	mercy	2	0	3	5	5	1
	worser	2	0	1	1	1	0

Term Frequency – actual data storage


- Inverted index saves only non-0 entries, not the whole matrix
 - Otherwise we would waste a lot of storage capacity
- Therefore not good at random lookups into the document column
 - Needs to iterate through the posting list to find the correct document
 - However, for scoring models $tf_{t,d}$ with 0 can be skipped



TF - Term Frequency

- $tf_{t,d}$ = how often does term t appear in document d
- Powerful starting point for many retrieval models
- Main point of our intuition at the beginning
- Using the raw frequency is not the best solution
 - Use relative frequencies
 - Dampen the values with logarithm

Term Frequency & Logarithm



- In long documents, a term may appear hundred of times.
- Retrieval experiments show that using the logarithm of the number of term occurrences is more effective than raw counts.
- Commonly used approach: apply logarithm

$$\log(1 + tf_{t,d})$$

Document Frequency

- df_t = in how many documents does term t appear in
 - Rare terms are more informative than frequent terms
 - Recall function words (and, or, the, ...)
 - Consider a term in the query that is rare in the collection
 - e.g., *Darmstadt* in a news corpora
 - A document containing this term is very likely to be relevant to the query *TU Darmstadt*
- We want a high weight for rare terms like *Darmstadt*.

IDF – Inverse Document Frequency

- A common way of defining the inverse document frequency of a term is as follows:

$$\text{idf}(t) = \log \frac{|D|}{df_t}$$

- df_t is an inverse measure of the “informativeness” of the term
- $df_t \leq |D|$
- Logarithm is used also for idf to “dampen” its effect.

$ D $	Total # of documents
df_t	# of Documents with $tf_t > 0$

TF-IDF

$$TF_IDF(q, d) = \sum_{t \in T_d \cap T_q} \underbrace{\log(1 + tf_{t,d})}_{\substack{\text{increases with the number of} \\ \text{occurrences within a document}}} * \underbrace{\log\left(\frac{|D|}{df_t}\right)}_{\substack{\text{increases with the rarity of} \\ \text{the term in the collection}}}$$

- A rare word (in the collection) appearing a lot in one document creates a high score
- Common words are downgraded

For more variations: <https://en.wikipedia.org/wiki/Tf-idf>

$\sum_{t \in T_d \cap T_q}$ Sum over all query terms, that are in the index

$tf_{t,d}$ Term frequency

$|D|$ Total # of documents

df_t # of Documents with $tf_{t,d} > 0$

TF-IDF – Usage

- Useful not only as a standalone model in document retrieval
- Weights used as a base for many other retrieval models
 - Example: Vector Space Model (VSM) works better with tf-idf weights
- Also useful as a generic word weighting mechanism for NLP
 - Task agnostic importance of a word in a document in a collection
 - Assign every word in a collection its tf-idf score

BM25

- Created 1994 by Robertson et al.
- Grounded in probabilistic retrieval
- In general, BM25 improves on TF-IDF results
- But only set as a default scoring in Lucene in 2015

Original paper: http://www.staff.city.ac.uk/~sb317/papers/robertson_walker_sigir94.pdf

TF-IDF vs BM25 in Lucene <https://opensourceconnections.com/blog/2015/10/16/bm25-the-next-generation-of-lucene-relevation/>

BM25 (as defined by Robertson et al. 2009)

$$BM25(q, d) = \sum_{t \in T_d \cap T_q} \frac{tf_{t,d}}{k_1((1-b) + b \frac{dl_d}{avgdl}) + tf_{t,d}} * \log \frac{|D| - df_t + 0.5}{df_t + 0.5}$$

- Simpler than the original formula
 - Over time it was shown that more complex parts not needed


Details (a lot of them): The Probabilistic Relevance Framework: BM25 and Beyond
http://www.staff.city.ac.uk/~sb317/papers/foundations_bm25_review.pdf

$\sum_{t \in T_d \cap T_q}$	Sum over all query terms, that are in the index
$tf_{t,d}$	Term frequency
dl_d	Document length
$avgdl$	Average document length in index
$ D $	Total # of documents
df_t	# of Documents with $tf_t > 0$
k_1, b	Hyperparameters

BM25 vs. TF-IDF

- Simple case of BM25 looks a lot like TF-IDF
- 1 main difference: BM25 tf component contains saturation function
 - Therefore works better in practice
- BM25 variants can be adapted to:
 - Incorporate additional reference information
 - Long(er) queries
 - multiple fields

BM25 vs. TF-IDF - Saturation



- **TF-IDF:** weight is always increasing (even with log)
- **BM25:** diminishing returns quickly = asymptotically approaches $k_1 + 1$

Note: we added (k_1+1) to the numerator to make $tf@1 = 1$, but it does not change the ranking because it is added to every term

Note: we assume the doc length = avgdl

BM25 vs. TF-IDF - Example

- Suppose your query is “machine learning”
- Suppose you have 2 documents with term counts:
 - doc1: learning 1024; machine 1
 - doc2: learning 16; machine 8
- TF-IDF: $\log(\text{tf}) * \log(|D|/\text{df})$
 - doc1: $11 * 7 + 1 * 10 = 87$
 - doc2: $5 * 7 + 4 * 10 = 75$
- BM25: $k_1 = 2$
 - doc1: $3 * 7 + 1 * 10 = 31$
 - doc2: $2.67 * 7 + 2.4 * 10 = 42.7$

Hyperparameters

- k_1, b are hyperparameters = they are set by us, the developers
- k_1 controls term frequency scaling
 - $k_1 = 0$ is binary model; k_1 large is raw term frequency
- b controls document length normalization
 - $b = 0$ is no length normalization; $b = 1$ is relative frequency (fully scale by document length)
- Common ranges: $0.5 < b < 0.8$ and $1.2 < k_1 < 2$

Summary: Part 1

- 1 We save statistics about terms in an inverted index
- 2 The statistics in the index can be access by a given term (query)
- 3 TF-IDF & BM25 use term and document frequencies to score a query & doc

Today

IR – Introduction, Evaluation

1 Introduction

- Inverted Index
- Search & Relevance
- TF-IDF & BM25

2 Evaluation

- Precision & Recall
- MRR & MAP
- nDCG

Evaluation

- We evaluate systems to observe concrete evidence for a hypothesis
 - Is our system better than the other one?
- IR systems are hard to evaluate
 - Ambiguity – what is relevant? In which context? Humans differ a lot ...
 - Collection size – explosion of query-document pairs
- Different types of result quality evaluation:
 - **Intrinsic:** Fixed set: same collection, query set & labels
 - **Extrinsic:** Observe behavior of users (in production system)*

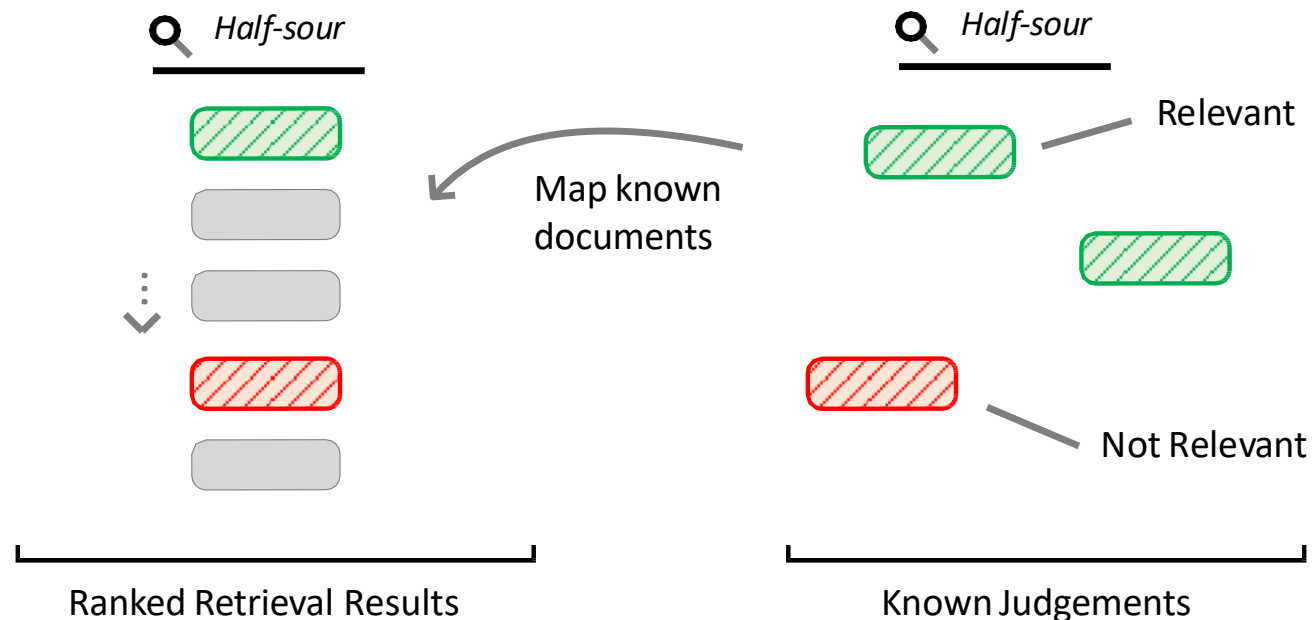
* Could also be a user study, beta version, etc...

The World of Evaluation

- Today we focus on evaluating the result quality of our own IR system
 - Does a document contain the answer for our query?
- Many other possibilities:
 - Efficiency
 - How fast can we index, return results for a query, how large becomes our index on disk?
 - Fairness, diversity, content quality, source credibility, effort, ...
 - Retrieval in the context of a larger goal
 - How many products, services do we sell through search
 - How well does our website integrate with Google, Bing, etc.. (SEO)
 - Optimizing a Blackbox


Extrinsic Evaluation Setup

- Quality of systems, that produce ranked list of documents
- Compared by a pool of judgements (does not necessarily cover the whole list)
 - Missing judgements are often considered as non-relevant



Comparing Systems

- We have multiple IR systems running on the same documents & same query
- How to compare them? Evaluation metrics to the rescue!



Today

IR – Introduction, Evaluation


1 Introduction

- Inverted Index
- Search & Relevance
- TF-IDF & BM25

2 Evaluation

- **Precision & Recall**
- MRR & MAP
- nDCG

Precision & Recall



How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

Precision/recall tradeoff

- You can increase recall (R) by returning more documents
 - Recall is a non-decreasing function of the number of documents retrieved
 - A system that returns all docs has 100% recall!
- The converse is also true: It's easy to get high precision (P) for very low recall
- Combined measure **F-score**:
- allows us to trade off precision against recall
- Mostly used measure: F1 or the harmonic mean of P and R

$$F_1 = 2 \times \frac{P \times R}{P + R}$$

Example for precision, recall, F1

	Relevant	Non-relevant	
Retrieved	20	40	60
Not retrieved	60	1,000,000	1,000,060
	80	1,000,040	1,000,120

$$P = \frac{20}{(20 + 40)} = \frac{1}{3}$$

$$R = \frac{20}{(20 + 60)} = \frac{1}{4}$$

$$F1 = 2 \times \frac{\frac{1}{3} \times \frac{1}{4}}{\frac{1}{3} + \frac{1}{4}} = \frac{2}{7}$$

	Relevant	Non-relevant
Retrieved	TP	FP
Not retrieved	FN	TN

$$P = \frac{TP}{TP + FP}$$

$$R = \frac{TP}{(TP + FN)}$$

$$F_1 = 2 \times \frac{P \times R}{P + R}$$

Today

IR – Introduction, Evaluation

1 Introduction

- Inverted Index
- Search & Relevance
- TF-IDF & BM25

2 Evaluation

- Precision & Recall
- **MRR & MAP**
- nDCG

Ranking List Evaluation Metrics

- Binary labels
 - **MRR**: Mean Reciprocal Rank
 - **MAP**: Mean Average Precision
- Graded labels
 - **nDCG**: normalized Discounted Cumulative Gain
- Typically we measure at a cutoff @k of the top retrieved documents
 - MAP, Recall: @100, @1000
 - Precision, MRR, nDCG: @5, @10, @20

Some nice explanations: <https://medium.com/swlh/rank-aware-recsys-evaluation-metrics-5191bba16832>

MRR: Mean Reciprocal Rank

Users look at results from the top; gets annoyed pretty fast; stops once they found the first relevant; doesn't care about the rest

$$MRR(Q) = \frac{1}{|Q|} * \sum_{q \in Q} \frac{1}{FirstRank(q)}$$

Mean over all queries


Reciprocal Rank

- MRR puts the focus on the first relevant document
- Applicable with sparse judgements or assuming users are satisfied with one relevant document

Q	$ Q $	$FirstRank(q)$
Query Set	Number of Queries	Returns the Rank of the first relevant document for 1 query

MRR: The Reciprocal Rank


- Reciprocal Rank: $\frac{1}{x}$
- Very strongly emphasis the first position



* x is plotted continuously, but in MRR x is discrete with the position in step size of 1

MRR: An Example

- Example for Reciprocal Rank:



MAP: Mean Average Precision

Users look at results closely, every time they find a new relevant document, they look at the full picture of what has been before

$$MAP(Q) = \frac{1}{|Q|} * \sum_{q \in Q} \frac{\sum_{i=1}^k P(q)_{@i} * rel(q)}{|rel(q)|}$$

Mean over all queries

Precision per relevant doc




Average Precision

- MAP squeezes complex evaluation into a single number
- Hard to interpret
- MAP corresponds to the area under the *Precision-Recall curve*

Q	$ Q $	$P(q)_{@}$	$rel(q)$	$ rel(q) $
Query Set	Number of Queries	Precision of query q after first i documents	Binary Relevance of doc at position i	Number of relevant documents

MAP: Mean Average Precision

- Example for Average Precision (2 relevant docs)
 - Mean is then calculated for multiple queries, for each system



Today

IR – Introduction, Evaluation

1 Introduction

- Inverted Index
- Search & Relevance
- TF-IDF & BM25

2 Evaluation

- Precision & Recall
- MRR & MAP
- **nDCG**

Graded Relevance

- Previous metrics all use binary relevance labels
 - Simple enough or too simple?
- Major problem: Of course there can be a difference in importance of relevance
 - Binary labels can not distinguish
- Graded relevance allows to assign different values of relevance
 - Can be floating point or fixed set of classes for manual annotation
 - Fixed set of classes for manual annotation
 - Floating point can be used when relevance inferred from logs

Common Graded TREC Relevance Labels

- [3] Perfectly relevant:** Document is dedicated to the query, it is worthy of being a top result in a search engine.
- [2] Highly relevant:** The content of this document provides substantial information on the query.
- [1] Relevant:** Document provides some information relevant to the query, which may be minimal.
- [0] Irrelevant:** Document does not provide any useful information about the query

nDCG: normalized Discounted Cumulative Gain

Users take for each document the relevance grade and position into account, normalize by best possible ranking per query

$$DCG(D) = \sum_{d \in D, i=1} \frac{rel(d)}{\log_2(i+1)}$$
$$nDCG(Q) = \frac{1}{|Q|} * \sum_{q \in Q} \frac{DCG(q)}{DCG(sorted(rel(q)))}$$

- nDCG compares actual results with maximum per query
- Relevance is graded
- nDCG@10 most commonly used in modern offline web search evaluation

Q	$ Q $	D	$rel(d)$	$rel(q)$	$sorted()$
Query Set	# of Queries	Single Doc. Result list	Relevance grade for single query-doc pair	List of all relevance grades for a query	Return graded documents by descending relevance

nDCG: A Closer Look

Discounted cumulative gain

$$DCG(D) = \sum_{d \in D, i=1} \frac{rel(d)}{\log_2(i+1)}$$

Gain (relevance value, commonly 0 -> 3)

Position Discounting

$$nDCG(Q) = \frac{1}{|Q|} * \sum_{q \in Q} \frac{DCG(q)}{DCG(sorted(rel(q)))}$$

Actual Results


Best possible sorting (ground truth)

Mean over all queries

Q	$ Q $	D	$rel(d)$	$rel(q)$	$sorted()$
Query Set	# of Queries	Single Doc. Result list	Relevance grade for single query-doc pair	List of all relevance grades for a query	Return graded documents by descending relevance

nDCG: Position Discounting




- Comparing the document position discount with reciprocal rank
 - Only for binary case $\text{rel}=1$
- nDCG discounts less than MRR



nDCG: Example

- Assuming two differently relevant docs (rel = 3 & 1)

- $$\text{Ideal DCG} = \frac{3}{\log_2(1+1)} + \frac{1}{\log_2(2+1)} = 3.63$$



Summary: Part 2

- 1 We compare systems with a set of query and document relevance labels
- 2 Binary metrics (MRR & MAP) are a solid foundation for evaluation
- 3 Graded relevance allows for more fine-grained metrics (nDCG)